

In []:

```
class FullyConnectedNet(object):
    """
    A fully-connected neural network with an arbitrary number of hidden layers,
    ReLU nonlinearities, and a softmax loss function. This will also implement
    dropout and batch normalization as options. For a network with L layers,
    the architecture will be

    {affine - [batch norm] - relu - [dropout]} x (L - 1) - affine - softmax

    where batch normalization and dropout are optional, and the {...} block is
    repeated L - 1 times.

    Similar to the TwoLayerNet above, learnable parameters are stored in the
    self.params dictionary and will be learned using the Solver class.
    """

    def __init__(self, hidden_dims, input_dim=3*32*32, num_classes=10,
                  dropout=1, use_batchnorm=False, reg=0.0,
                  weight_scale=1e-2, dtype=np.float32, seed=None):
        """
        Initialize a new FullyConnectedNet.

        Inputs:
        - hidden_dims: A list of integers giving the size of each hidden layer.
        - input_dim: An integer giving the size of the input.
        - num_classes: An integer giving the number of classes to classify.
        - dropout: Scalar between 0 and 1 giving dropout strength. If dropout=1 then
          the network should not use dropout at all.
        - use_batchnorm: Whether or not the network should use batch normalization.
        - reg: Scalar giving L2 regularization strength.
        - weight_scale: Scalar giving the standard deviation for random
          initialization of the weights.
        - dtype: A numpy datatype object; all computations will be performed using
          this datatype. float32 is faster but less accurate, so you should use
          float64 for numeric gradient checking.
        - seed: If not None, then pass this random seed to the dropout layers. This
          will make the dropout layers deterministic so we can gradient check the
          model.
        """
        self.use_batchnorm = use_batchnorm
        self.use_dropout = dropout < 1
        self.reg = reg
        self.num_layers = 1 + len(hidden_dims)
        self.dtype = dtype
        self.params = {}

        # ===== #
        # YOUR CODE HERE:
        # Initialize all parameters of the network in the self.params dictionary.
        # The weights and biases of layer 1 are W1 and b1; and in general the
        # weights and biases of layer i are Wi and bi. The
        # biases are initialized to zero and the weights are initialized
        # so that each parameter has mean 0 and standard deviation weight_scale.
        #
        # BATCHNORM: Initialize the gammas of each layer to 1 and the beta
        # parameters to zero. The gamma and beta parameters for layer 1 should
        # be self.params['gamma1'] and self.params['beta1']. For layer 2, they
        # should be gamma2 and beta2, etc. Only use batchnorm if self.use_batchnorm
        # is true and DO NOT do batch normalize the output scores.
        # ===== #

        #we will initialize them based on layers:
        for i in np.arange(1, self.num_layers + 1): #for all layers

            if i == 1: #for the layer 1

                self.params['b1'] = np.zeros(hidden_dims[i - 1]) #based on the size of the first hidden layer
                #the weights will have mean 0 and standard deviation weight_scale
                self.params['W1'] = np.random.normal(0, weight_scale, (input_dim, hidden_dims[i - 1]))

                #if the network uses batch normalization
                if self.use_batchnorm == True:
                    self.params['gamma1'] = np.ones(hidden_dims[i - 1]) #the gammas should be 1
                    self.params['beta1'] = np.zeros(hidden_dims[i - 1]) #the betas should be zero

            elif i == self.num_layers: #for the last layer

                #the size after the last layer will be the number of classes
                self.params['b' + str(i)] = np.zeros(num_classes)
                self.params['W' + str(i)] = np.random.normal(0, weight_scale, (hidden_dims[i - 2], num_classes))

            else: #for the layer except the first and the last

                #based on the size of the ith hidden layer
```

```

        self.params['b' + str(i)] = np.zeros(hidden_dims[i - 1])
        self.params['W' + str(i)] = np.random.normal(0, weight_scale, (hidden_dims[i - 2], hidden_dims[i - 1]))

        #if the network uses batch normalization
        if self.use_batchnorm == True:
            self.params['gamma' + str(i)] = np.ones(hidden_dims[i - 1]) #the gammas should be 1
            self.params['beta' + str(i)] = np.zeros(hidden_dims[i - 1]) #the betas should be zero

# ===== #
# END YOUR CODE HERE
# ===== #

# When using dropout we need to pass a dropout_param dictionary to each
# dropout layer so that the layer knows the dropout probability and the mode
# (train / test). You can pass the same dropout_param to each dropout layer.
self.dropout_param = {}
if self.use_dropout:
    self.dropout_param = {'mode': 'train', 'p': dropout}
if seed is not None:
    self.dropout_param['seed'] = seed

# With batch normalization we need to keep track of running means and
# variances, so we need to pass a special bn_param object to each batch
# normalization layer. You should pass self.bn_params[0] to the forward pass
# of the first batch normalization layer, self.bn_params[1] to the forward
# pass of the second batch normalization layer, etc.
self.bn_params = []
if self.use_batchnorm:
    self.bn_params = [{'mode': 'train'} for i in np.arange(self.num_layers - 1)]

# Cast all parameters to the correct datatype
for k, v in self.params.items():
    self.params[k] = v.astype(dtype)

def loss(self, X, y=None):
    """
    Compute loss and gradient for the fully-connected net.

    Input / output: Same as TwoLayerNet above.
    """
    X = X.astype(self.dtype)
    mode = 'test' if y is None else 'train'

    # Set train/test mode for batchnorm params and dropout param since they
    # behave differently during training and testing.
    if self.dropout_param is not None:
        self.dropout_param['mode'] = mode
    if self.use_batchnorm:
        for bn_param in self.bn_params:
            bn_param['mode'] = mode

    scores = None

# ===== #
# YOUR CODE HERE:
# Implement the forward pass of the FC net and store the output
# scores as the variable "scores".
#
# BATCHNORM: If self.use_batchnorm is true, insert a bathnorm layer
# between the affine_forward and relu_forward layers. You may
# also write an affine_batchnorm_relu() function in layer_utils.py.
#
# DROPOUT: If dropout is non-zero, insert a dropout layer after
# every ReLU layer.
# ===== #

    out_list = []
    cache_list = []
    dropout_list = []
    X = X.reshape([X.shape[0], -1])

    for i in np.arange(1, self.num_layers + 1):

        #we will consider three instances:

        if i == 1: #for the layer 1

            #if the network uses batch normalization
            if self.use_batchnorm:
                #affine -> batchnorm -> relu
                out1, cache1 = affine_batchnorm_relu_forward(X, self.params['W' + str(i)], self.params['b' + str(i)],
                    self.params['gamma' + str(i)], self.params['beta' + str(i)], self.bn_params[i]

            else:
                #affine -> relu

```

```

        out1, cache1 = affine_relu_forward(X, self.params['W' + str(i)], self.params['b1'])

    cache_list.append(cache1)

    #If dropout is non-zero
    if self.use_dropout:

        #insert a dropout after ReLU layer
        out1, cached = dropout_forward(out1, self.dropout_param)
        dropcache_list.append(cached)

    out_list.append(out1)

elif i == self.num_layers: #for the last layer

    #do the last layer affine forward
    scores, cachescores = affine_forward(out_list[i - 2], self.params['W' + str(i)], self.params['b' + str(i)])

    cache_list.append(cachescores)

else: #for the layers except the first and the last

    #if the network uses batch normalization
    if self.use_batchnorm:
        #affine -> batchnorm -> relu
        outnow, cachenow = affine_batchnorm_relu_forward(out_list[i - 2], self.params['W' + str(i)], self.params['gamma' + str(i)], self.params['beta' + str(i)], self.params['b' + str(i)])
    else:
        #affine -> relu
        outnow, cachenow = affine_relu_forward(out_list[i - 2], self.params['W' + str(i)], self.params['b' + str(i)])

    cache_list.append(cachenow)

    #If dropout is non-zero
    if self.use_dropout:

        #insert a dropout after ReLU layer
        outnow, cached = dropout_forward(outnow, self.dropout_param)
        dropcache_list.append(cached)

    out_list.append(outnow)

# ===== #
# END YOUR CODE HERE
# ===== #

# If test mode return early
if mode == 'test':
    return scores

loss, grads = 0.0, {}
# ===== #
# YOUR CODE HERE:
# Implement the backwards pass of the FC net and store the gradients
# in the grads dict, so that grads[k] is the gradient of self.params[k]
# Be sure your L2 regularization includes a 0.5 factor.
#
# BATCHNORM: Incorporate the backward pass of the batchnorm.
#
# DROPOUT: Incorporate the backward pass of dropout.
# ===== #

loss, dx = softmax_loss(scores, y) #calculate loss and dx by given softmax_loss

for j in np.arange(self.num_layers, 0, -1): #loop from the last layer to the first layer
    #L2 regularization by multiplying the regularization loss by 0.5(with factor reg)
    loss += (self.reg) * (np.sum(self.params['W' + str(j)] ** 2)) * 0.5

    if j == self.num_layers: #from the last layer
        #backward affine of the second layer
        currentdx, grads['W' + str(j)], grads['b' + str(j)] = affine_backward(dx, cache_list[j - 1])

    else: #for other layers:

        #If dropout is non-zero
        if self.use_dropout:
            currentdx = dropout_backward(currentdx, dropcache_list[j - 1])

        #if the network uses batch normalization
        if self.use_batchnorm:
            currentdx, grads['W' + str(j)], grads['b' + str(j)], grads['gamma' + str(j)], grads['beta' + str(j)] = affine_batchnorm_relu_backward(currentdx, cache_list[j - 1])
        else:
            currentdx, grads['W' + str(j)], grads['b' + str(j)] = affine_relu_backward(currentdx, cache_list[j - 1])

```

```
    #update the gradients of W
    grads['W' + str(j)] += (self.reg) * self.params['W' + str(j)]

# ===== #
# END YOUR CODE HERE
# ===== #

    return loss, grads
```